

First 120 Days: Earn the Right to Standardize

A discovery-led entry plan for enterprise AI strategy, data governance, assurance, business partnership, and measurable adoption.

Entry principles

Respect existing authority Do not centralize decisions that belong to engineering, quality, security, legal, privacy, regulatory, or business owners.	Start with live work Test governance in representative use cases instead of designing only from policy abstractions.	Measure the mechanism Track evidence gaps, review latency, exceptions, overrides, adoption, and outcomes before declaring maturity.
---	--	---

Outcome by day 120

A visible AI portfolio; a co-owned assurance-case standard; consequence-based operating postures; clarified decision rights; representative use cases that prove or challenge the mechanism; an AI business-partner playbook; and an executive scorecard with defined units, owners, sources, and decision use.

Stakeholder topology

Stakeholder	What I need to learn	What they should retain or gain
CTO / executive sponsors	Strategic priorities, investment logic, risk appetite, decision bottlenecks.	Clear portfolio choices, evidence, exception visibility, and value measures.
Business and AI partners	Demand patterns, workflow pain, adoption burden, ownership gaps.	Faster qualified intake, reusable controls, and outcome accountability.
Data leaders / stewards	Domains, lineage, quality, access, catalog, retention, ownership.	Governance tied to priority AI use and explicit remediation value.
Technology / platform teams	Architecture, model/tool landscape, monitoring, release, recovery.	Clear evidence contract, approved patterns, and exception path.
Independent authorities	Security, privacy, legal, quality, regulatory, records, export-control requirements.	Preserved authority, earlier engagement, auditable decisions, fewer late surprises.

Days 0-30 - Read the system

Workstream	Actions	Evidence / output
Portfolio	Inventory active, planned, paused, and informal AI initiatives; connect each to a workflow, owner, decision, and business outcome.	Portfolio map with source, owner, status, consequence, dependencies, and unanswered questions.
Governance flow	Trace how an idea becomes a funded pilot, production capability, monitored service, exception, or stop decision.	Current decision flow; review latency; duplicate approvals; missing owners; recurring exceptions.
Data	Map priority domains, lineage/catalog practices, quality issues, access, retention, stewardship, and cross-region constraints.	Data dependency map and prioritized gaps linked to business use cases.
Authority	Interview business, engineering, quality, security, privacy, legal, regulatory, data, and platform leaders.	Decision-rights map: recommend, approve, validate, operate, monitor, stop, appeal.
Adoption	Observe how users receive, challenge, correct, ignore, or work around AI outputs.	Adoption burden and human-authority evidence, not satisfaction theater.

Days 31-60 - Define the contract

Deliverable	Design question	Acceptance evidence
Assurance-case standard	What must every use case make explicit regardless of technology?	Six evidence dimensions with owners, minimum records, review triggers, and change controls.
Operating postures	How should evidence burden vary by consequence and human authority?	Explore, assist, control, and independent-review paths with clear entry/exit conditions.
Exception model	How are deviations documented, time-bounded, monitored, escalated, and retired?	Named approver, rationale, compensating control, review date, evidence, closure.
Portfolio scorecard	Which measures change funding, scale, hold, remediation, or stop decisions?	Defined units, source systems, owners, cadence, and decision use; no invented targets.
Pilot set	Which use cases best stress different data, authority, and consequence patterns?	Representative cases with willing owners and a real operating decision to improve.

Days 61-90 - Prove the mechanism

- Run assurance-case workshops inside live delivery work with business, data, technical, and independent authorities together.
- Measure time to evidence, review latency, exception aging, rework, unresolved authority, human overrides, and adoption behavior.
- Test whether the standard reduces ambiguity and late surprises without creating clerical work that does not change a decision.
- Convert successful evidence and control patterns into reusable templates, platform guardrails, test suites, and data requirements.
- Document where the model fails, where authority is unclear, and which controls should be retired, moved earlier, automated, or strengthened.

Days 91-120 - Scale the learning

- Launch an executive portfolio cadence that distinguishes exploration, qualified pilots, controlled production, remediation, exceptions, and stopped work.
- Equip AI business partners with an intake and discovery playbook that begins with outcome, workflow, owner, human authority, evidence, adoption, and consequence.
- Publish the assurance-case standard, operating postures, decision-rights model, exception protocol, and control-pattern library.
- Set the next portfolio using evidence from the pilots rather than enthusiasm or generic maturity scoring.
- Establish the regulatory/standards watch and change process for NIST, NRC, ISO/IEC, EU, privacy, cybersecurity, and relevant customer expectations.

Day-120 executive readout

What changed	What we learned	What to decide
Portfolio visibility, decision clarity, evidence reuse, data priorities, adoption ownership.	Where governance accelerates work, where it adds burden, where risk or authority remains unclear.	Fund, scale, remediate, standardize, pause, stop, or escalate - with evidence attached.

Risks I would manage explicitly

Risk	Response
Governance becomes a centralized queue	Federated decision rights, reusable evidence, clear thresholds, and exception handling.
Business partners become idea collectors	Measure qualified demand, adoption, workflow outcomes, and benefit realization.
Data governance remains detached	Prioritize lineage, quality, and stewardship through funded use cases and consequence.
Metrics create false precision	Define units, sources, owners, and decisions first; set targets only after baseline validation.
AI authority is ambiguous	Make human decision, override, appeal, and stop rights explicit before production.