

Make assurance portable.

Interview thesis for Director, AI Strategic Governance: use a reusable assurance case to connect enterprise AI strategy, data governance, risk, business partnership, adoption, and learning.

Actual mandate

Turn expanding AI demand into a governed portfolio of adopted capabilities that create measurable operating value while preserving evidence, authority, cybersecurity, privacy, quality, and regulatory trust.

Company moment

- Westinghouse publicly positions AI across reactor design, licensing, procurement, manufacturing, construction, and operating-fleet improvement.
- Its growth agenda increases the premium on standardization, repeatability, secure data, and execution credibility.
- The role reports to the CTO and combines AI strategy, data strategy, governance, enterprise partnership, and executive/regulator communication.

Central tension

**Move fast enough to create enterprise leverage;
move visibly enough to preserve nuclear-grade
trust and human authority.**

Candidate thesis

Every AI use case should carry a portable assurance case that makes its purpose, data, behavior, authority, controls, and outcomes explicit.

Three contradictions to manage

Speed vs evidence

Pilots compound hidden debt when assurance begins after delivery.

Control vs reuse

Bespoke review protects locally but blocks enterprise learning.

Centralization vs authority

Central coordination cannot replace qualified functional decisions.

What I would build

A practical AI governance system that travels with the work. Each use case would carry its purpose, data and model evidence, human decision authority, applicable controls, and operating results. Evidence requirements would increase with consequence: bounded experiments could move quickly, while production and high-consequence uses would require stronger validation, monitoring, recovery, and independent review. The result would be faster decisions, clearer accountability, reusable controls, and a portfolio that learns from evidence instead of repeatedly starting from zero.

Research basis: Westinghouse AI and Google Cloud announcements, current role posting, public leadership and growth communications; NIST AI RMF/GenAI Profile, NRC AI strategy, ISO/IEC 42001, and EU AI Act. Company-specific operating details are hypotheses for discovery.

OPERATING MODEL

The AI Assurance Lattice

A stable six-dimension evidence contract surrounds each use case. Evidence depth and review authority change by consequence; the conceptual grammar stays portable.

1. Mission & boundary Business outcome, prohibited decisions, accountable owner, affected workflow.	2. Data pedigree Source, ownership, lineage, quality, access, retention, representativeness.	3. Model evidence Evaluation, uncertainty, robustness, drift, failure modes, change control.
4. Human authority Named decision owner, review, override, appeal, escalation, stop rights.	5. Security & compliance Cybersecurity, privacy, legal, export control, quality, records, regulation.	6. Outcome telemetry Adoption, workflow results, overrides, incidents, economics, learning.

Consequence-based operating postures

Posture	Use	Minimum governance behavior
Explore	Bounded sandbox and discovery	Approved data, prohibited-use boundary, accountable sponsor, no operational reliance.
Assist	Human-authoritative workflow support	Grounding/evaluation evidence, visible sources, trained user, clear rejection and correction path.
Control	Production workflow with monitoring	Versioned validation, security/privacy controls, monitoring, rollback, incident and exception records.
Independent review	High-consequence or regulated decision support	Qualified independent authority, formal assurance evidence, challenge testing, strict change and records controls.

Decision architecture

Central AI governance Portfolio, reusable patterns, exception framework, executive transparency, external scanning.	Federated authority Business owners own outcomes; data stewards own data; technical owners own operation; independent functions retain approval and stop rights.
---	--

EVIDENCE MAP

Why this operating problem fits my record

Mandate need	Verified evidence	Transfer logic
High-consequence governance	Compunetics quality systems and high-reliability engineering; ZeusVu regulated operations with 100% safety record and no FAA incidents.	Mechanism discipline, evidence, controls, root cause, independent authority, and respect for consequence.
Enterprise AI strategy	DudeWorth AI readiness and opportunity frameworks; agentic workflows with retrieval, human judgment, evaluation, escalation, and governance rails.	Connect AI ambition to value, feasibility, data, authority, adoption, and measurable work.
Adoption and operating cadence	Amazon launch and 24/7 operations; Vape-Jet cross-functional AI-first operating transformation.	Translate strategy into routines, ownership, training, dashboards, escalation, and sustained use.
Data/workflow integration	Vape-Jet ERP, CRM, Jira, Slack, telemetry, quality, support, engineering and machine-fleet signals.	Governance succeeds when evidence and decision rights live inside work rather than separate from it.
Executive/business partnership	Director-level engineering/operations, regional operations, board leadership, IEEE ecosystem programming.	Communicate across executives, technical specialists, operations, customers, and public/private stakeholders.

Balanced executive scorecard

Portfolio economics Funded outcomes, benefit realization, stopped or reused work.	Delivery flow Time to qualified pilot, review latency, exception aging, rework.	Assurance quality Lineage/evidence completeness, validation defects, incidents, recovery.
Human authority Override patterns, appeal/stop paths, unresolved ownership.	Adoption Workflow use, training, repeat use, correction behavior, owner engagement.	Learning speed Reusable patterns, retired controls, updated standards, monitored drift.

No numeric targets are proposed without a validated Westinghouse baseline. The first governance scorecard should define units, sources, owners, and decision use before setting thresholds.

Strongest reasonable concern, answered affirmatively

My career is adjacent to nuclear rather than inside a nuclear OEM. The day-one value is not false domain equivalence; it is the combination of high-reliability technical operations, quality systems, regulated autonomy, AI workflow/governance design, scaled adoption, and scientific grounding. I would use that background to strengthen the operating mechanism while qualified Westinghouse experts retain nuclear decisions.

FIRST 120 DAYS

Earn the right to standardize

Phase	Purpose	Deliverables and evidence
0-30 Read	Understand the current AI/data system, pressure, authority, and friction.	Portfolio map; decision-rights map; stakeholder interviews; current review/value flow; source-grounded regulatory watchlist.
31-60 Define	Co-design the portable assurance contract and consequence postures.	Assurance-case standard; data and evidence ownership; exception model; pilot selection; draft scorecard definitions.
61-90 Prove	Install the mechanism inside representative delivery work.	Live assurance cases; cycle-time and evidence-gap findings; redundant-control retirement; adoption and authority feedback.
91-120 Scale	Launch the portfolio cadence and reusable patterns.	Executive portfolio review; AI business-partner playbook; control-pattern library; next-wave investment/stop decisions.

Executive interview questions

1. Where does the current governance system create confidence, and where does it create avoidable cycle time?
2. Which AI decisions require global consistency, and which must remain with engineering, quality, security, legal, privacy, regulatory, regional, or business authorities?
3. What evidence changes an executive decision from explore to fund, from pilot to scale, or from continue to stop?
4. How will the AI business partner team be measured across demand quality, adoption, outcomes, and control improvement?
5. Which external expectations will materially change the minimum assurance case in the next 18 months?
6. What would make this function trusted by both innovators and independent control authorities?

90-second closing thesis

Westinghouse's opportunity is not simply to approve more AI or write more policy. It is to create a repeatable way for each use case to prove what it does, what data it uses, how it behaves, who retains authority, which controls apply, and whether it improves the work. When that assurance case is portable, governance becomes a source of velocity and institutional learning. That is the system I would help build - with nuclear expertise retaining the authority it has earned.

Supporting campaign

Candidate vision: <https://russelldudek.github.io/Laveer/>
Includes role-aligned resume, cover letter, full 120-day entry plan, and proposed AI Assurance Case Standard.